

Name: VINAYAK AABUJ,

TUSHAR VERMA

Project Report: Q-Learning for Maze Navigation

Problem Statement

The objective of this project is to implement a Q-Learning agent capable of navigating a grid-based maze to find the shortest path from a given start point to a goal while avoiding obstacles. This task showcases the practical application of reinforcement learning in solving pathfinding problems in dynamic and constrained environments.

Real-World Relevance

1. **Autonomous Navigation:** The principles used in this project are applicable in designing navigation systems for robots or drones in unknown terrains.
2. **Logistics and Supply Chain:** Optimized route planning in constrained warehouse spaces.
3. **Gaming and Simulation:** AI agents in games and simulations can benefit from similar reinforcement learning techniques.

Approach and Methodology

Environment Design

- The maze is modeled as a 2D grid where:
 - 0 represents free cells.
 - 1 represents obstacles.
 - The agent can start from a specific cell and must reach the goal cell.

State Representation

- Each state corresponds to the agent's current grid cell.

Action Space

- The agent can choose from four possible actions: up, down, left, right.

Reward System

- +100 for reaching the goal.
- -10 for hitting obstacles.
- -1 for each step to encourage shorter paths.

Q-Learning Algorithm

1. Initialization:

- Initialize a Q-table with zeros for all state-action pairs.

2. Policy:

- Use an ϵ -greedy policy to balance exploration (choosing random actions) and exploitation (choosing actions based on the Q-table).

3. Update Rule:

- The Q-value is updated using the formula:

$$Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma Q(s',a) - Q(s,a)]$$

- where:

- α is the learning rate.
- γ is the discount factor.
- r is the reward for the action.
- s' is the next state.

4. Training:

- The agent learns by iteratively updating the Q-table over multiple episodes.

Visualization

- A plot of the maze with the agent's optimal path is generated to visualize the learned policy.

Results and Observations

Training Results

- After training for 1000 episodes:
 - The agent successfully learned an optimal path to the goal.
 - The Q-table values converged, indicating a stable policy.

Path Visualization

- The agent's learned path avoids obstacles and reaches the goal in the shortest time possible, as observed in the generated plots.

Performance Analysis

- The agent performed well in mazes with clear pathways.

- For mazes with narrow corridors or high obstacle density, convergence required more episodes.

Challenges Faced

1. Sparse Rewards:

- The agent initially struggled to find the goal due to sparse positive rewards in the environment.

2. Exploration-Exploitation Tradeoff:

- Balancing exploration (to discover new paths) and exploitation (to use known good paths) was challenging.

3. Hyperparameter Tuning:

- Parameters like the learning rate (α), discount factor (γ), and exploration rate (ϵ) required careful tuning for optimal performance.

Potential Improvements

1. Dynamic Exploration Rate:

- Gradually reducing ϵ over episodes could allow more exploration early on and more exploitation later.

2. Reward Shaping:

- Providing intermediate rewards for moving closer to the goal could accelerate learning.

3. Complex Environments:

- Testing the algorithm on larger mazes with more complex obstacle patterns.

4. Deep Reinforcement Learning:

- Extending the project to use neural networks to approximate the Q-values for environments with larger state spaces.